

CPE 213 Data Models

(a.k.a. Data Modeling and Visualization)

Lecture 2: R Tutorial

Asst. Prof. Dr. Santitham Prom-on

Department of Computer Engineering, Faculty of Engineering
King Mongkut's University of Technology Thonburi

Lecture 2 Overview

- R basic
- Data manipulation with tidyr and dplyr

<http://fastdata.in.th/data-model-2021/>

R Basics

Section 1

What is R?

- R is a scripting language for data manipulation and analysis
- R is originated from S language at Bell Laboratories (AT&T)
 - R and S were created with purposes of providing an interactive environment without requiring programming skills
- Over the years, through a lot of academic and commercial contributions, numerous packages were developed with various purposes and better qualities.

IEEE Ranking of Popular Programming Language 2016


Language Ranking: IEEE Spectrum

| Rank | Language | Type | Score |
|------|------------|---|-------|
| 1 | Python |    | 100.0 |
| 2 | Java |    | 95.3 |
| 3 | C |    | 94.6 |
| 4 | C++ |    | 87.0 |
| 5 | JavaScript |  | 79.5 |
| 6 | R |  | 78.6 |
| 7 | Arduino |  | 73.2 |
| 8 | Go |   | 73.1 |
| 9 | Swift |   | 70.5 |
| 10 | Matlab |  | 68.4 |

<https://spectrum.ieee.org/static/interactive-the-top-programming-languages-2020>

KDnuggets

Most Requested programming languages for Data Science in 2019



Why use R?

- Free
- Easy to get the job done
- Clear, more compact code
- Easier transition to distributed computing
- Great supporting community
- Object-oriented programming
- Functional programming

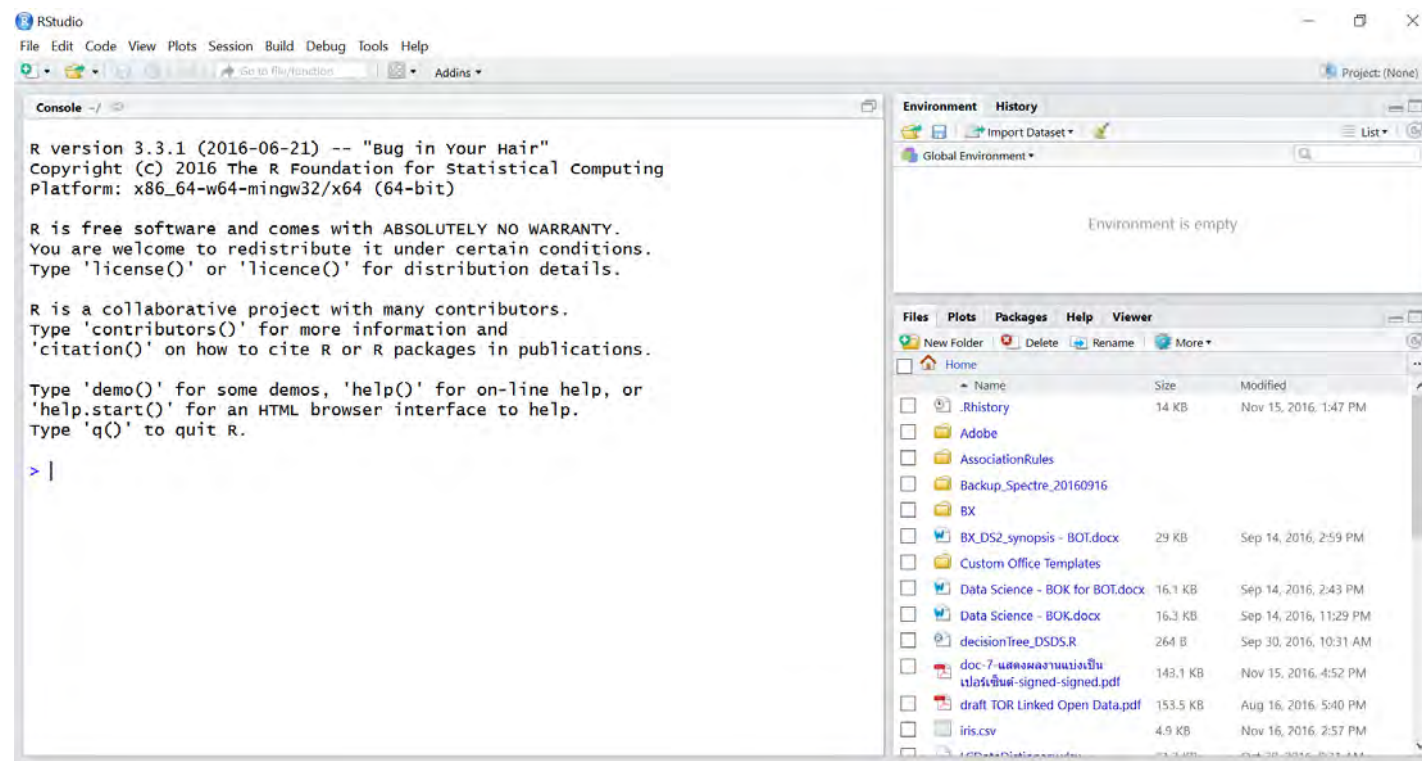
Learning R

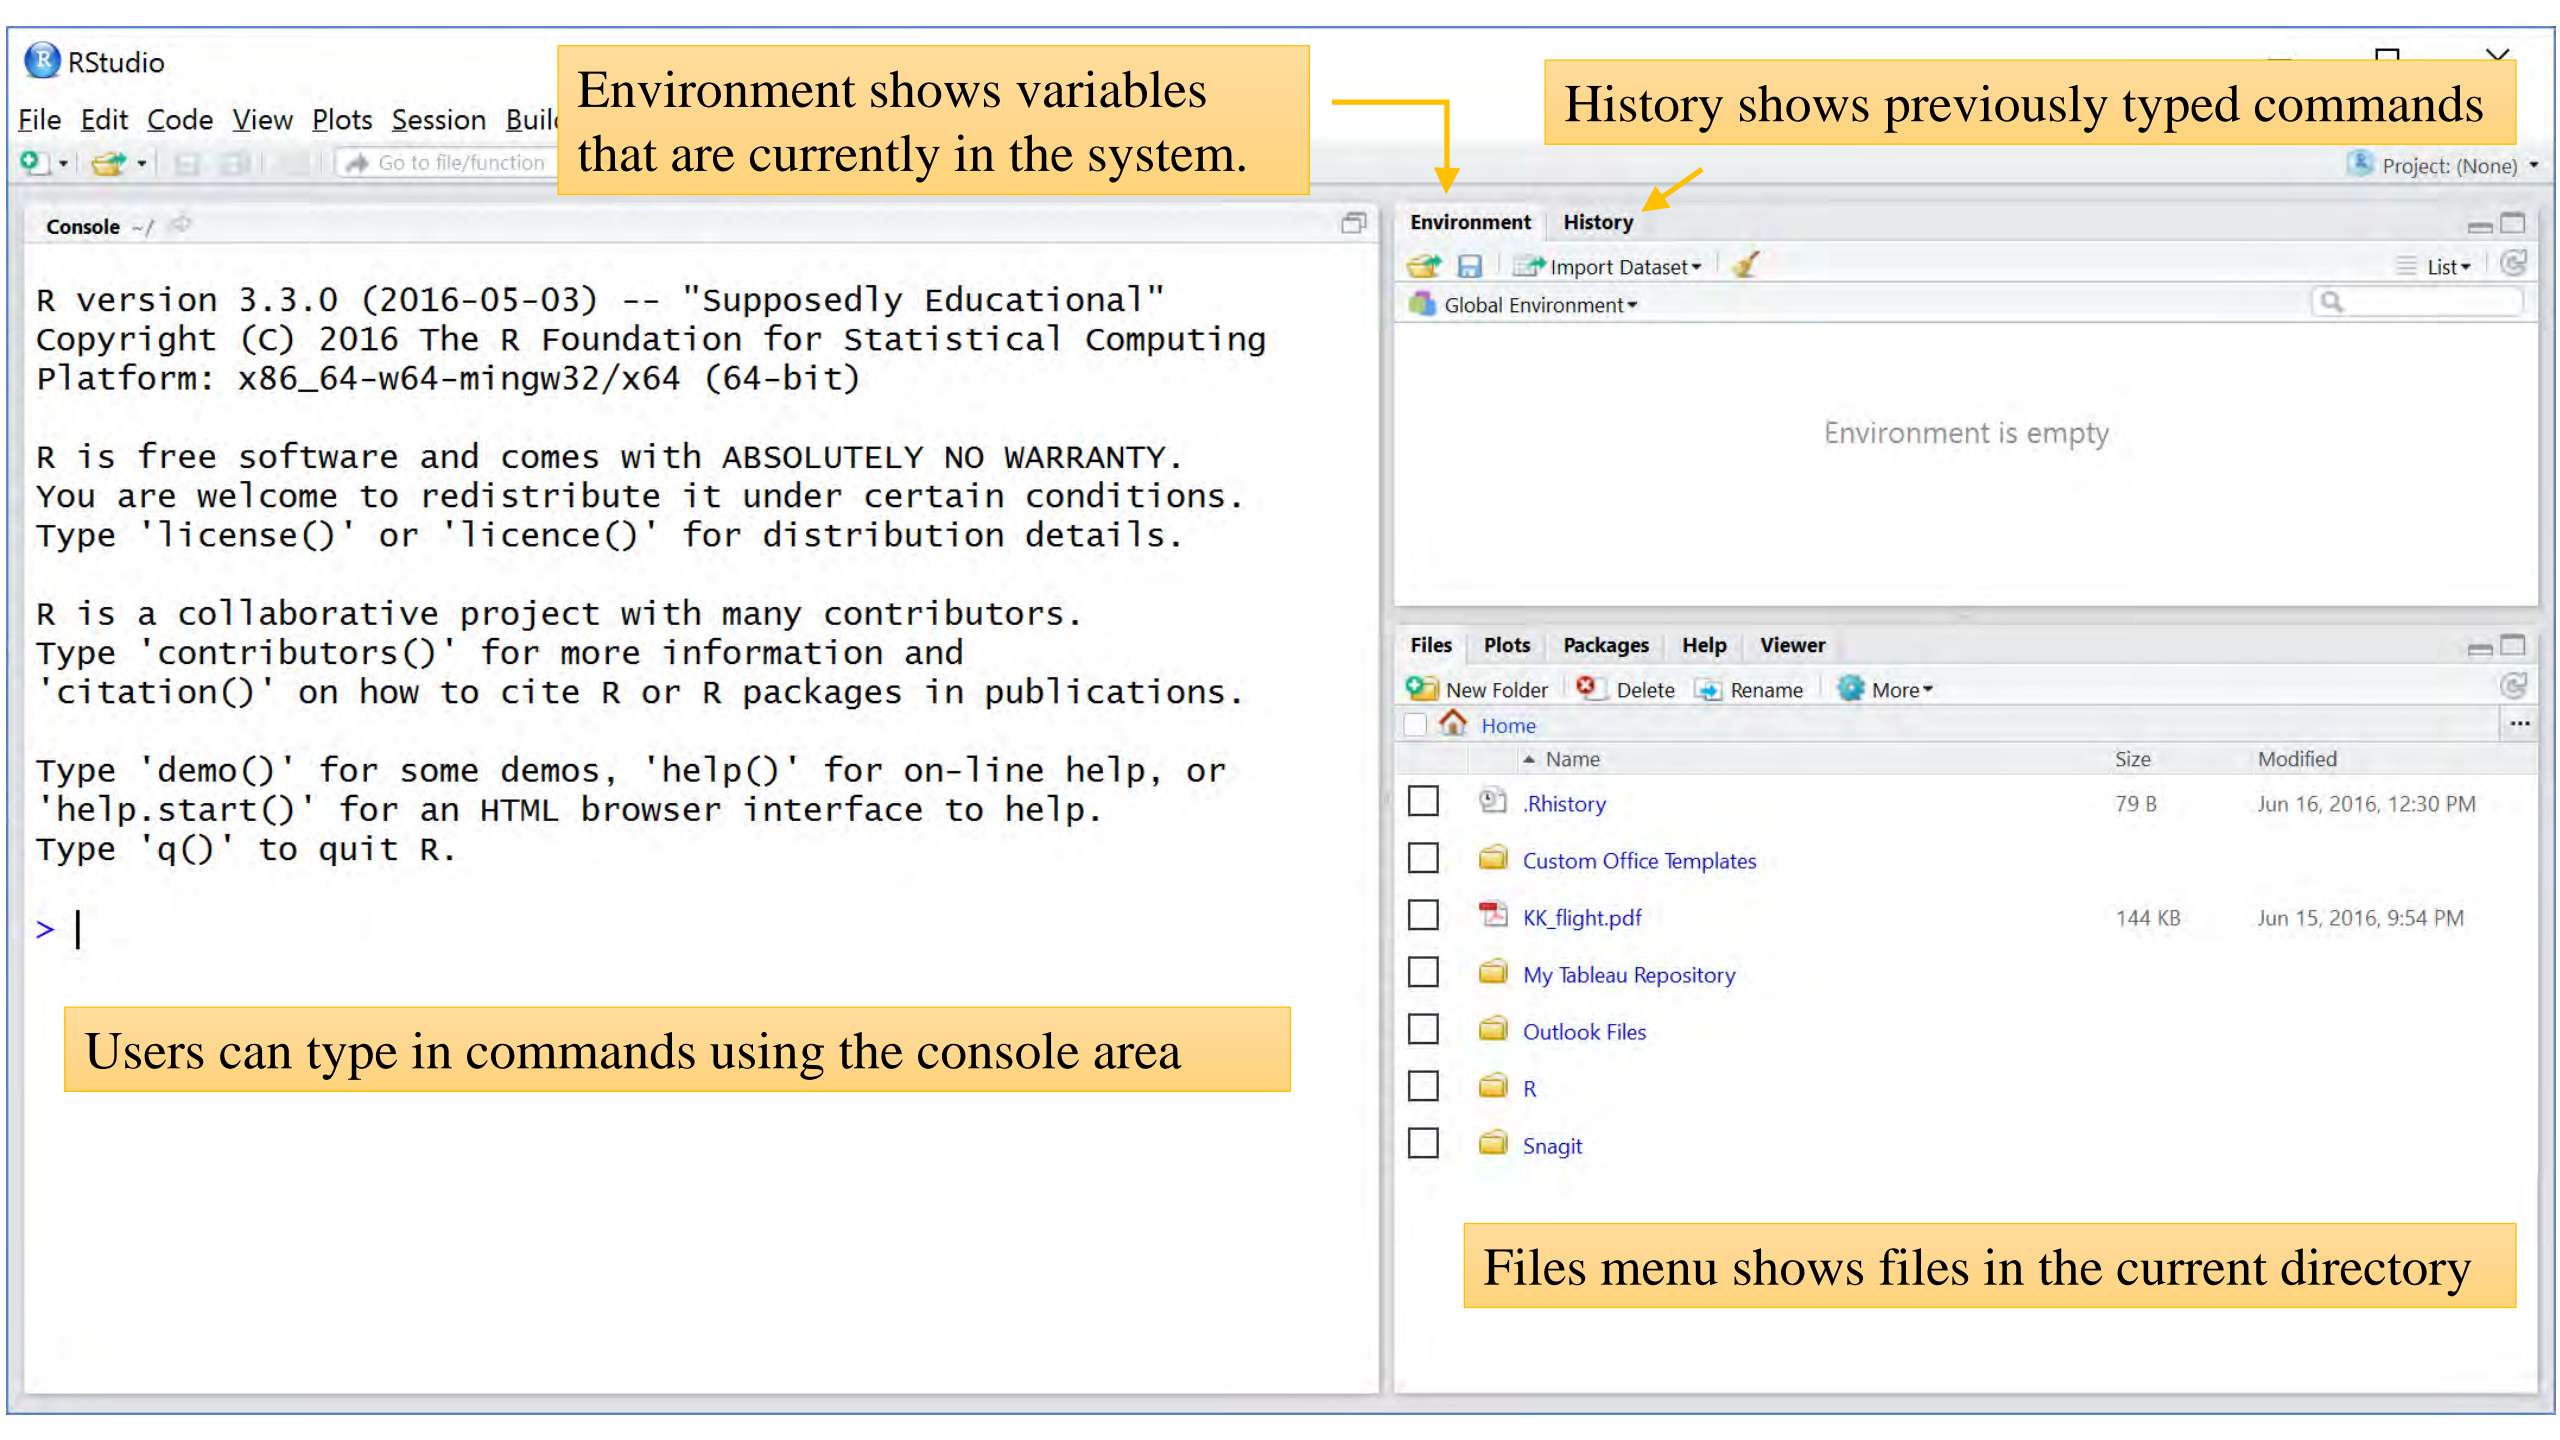
- Learning R is not like learning C or Java languages
- R is a dynamic typing (no declaration required)
- R is functional (use function to run)
- R is data oriented (utilize data structure to do the work)
- R is learned through an interactive exploration of how to interact with the environment and data

Rstudio (via www.rstudio.com)



- RStudio is an IDE for R
- RStudio makes R easier to use
 - Console
 - Syntax highlight editor
 - Plotting
 - History
 - Environment
 - Help
 - Debugging





Environment shows variables that are currently in the system.

History shows previously typed commands

Console ~/

R version 3.3.0 (2016-05-03) -- "Supposedly Educational"
Copyright (C) 2016 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> |

Users can type in commands using the console area

Files | Plots | Packages | Help | Viewer

New Folder | Delete | Rename | More

Home

| | Name | Size | Modified |
|--------------------------|-------------------------|--------|------------------------|
| <input type="checkbox"/> | .Rhistory | 79 B | Jun 16, 2016, 12:30 PM |
| <input type="checkbox"/> | Custom Office Templates | | |
| <input type="checkbox"/> | KK_flight.pdf | 144 KB | Jun 15, 2016, 9:54 PM |
| <input type="checkbox"/> | My Tableau Repository | | |
| <input type="checkbox"/> | Outlook Files | | |
| <input type="checkbox"/> | R | | |
| <input type="checkbox"/> | Snagit | | |

Files menu shows files in the current directory

Try this...

Type these commands and guess their functions

```
1+1
```

```
mean(runif(100))
```

```
1:10
```

```
2**10
```

```
abs(rnorm(10))
```

Variable assignment

```
x <- c(1, 2, 4)
```

- This function creates a vector of 1, 2 and 4
- There are no fixed types associated with variables
- `c()` is a function that concatenates three one-element vector
- The assignment operator in R are `<-` and `=`
- `<-` is a preferred assignment operator while `=` is used for parameter assignment
- Try this `q <- c(x, x, 8)`

Variable printing and subsetting

```
> q
[1] 1 2 4 1 2 4 8
> q[1]
[1] 1
> q[1:4]
[1] 1 2 4 1
> q[-2]
[1] 1 4 1 2 4 8
> q[-1:-4]
[1] 2 4 8
> q[c(2,5)]
[1] 2 2
```

Print a variable `q`

Select an individual element of a vector, index 1

Subsetting range

Deselect

Deselect range

Selective subsetting

Calling functions

```
y <- mean(x)
```

```
y1 <- sd(x)
```

- The above code calculate mean and standard deviation of `x` and store them to `y` and `y1` respectively

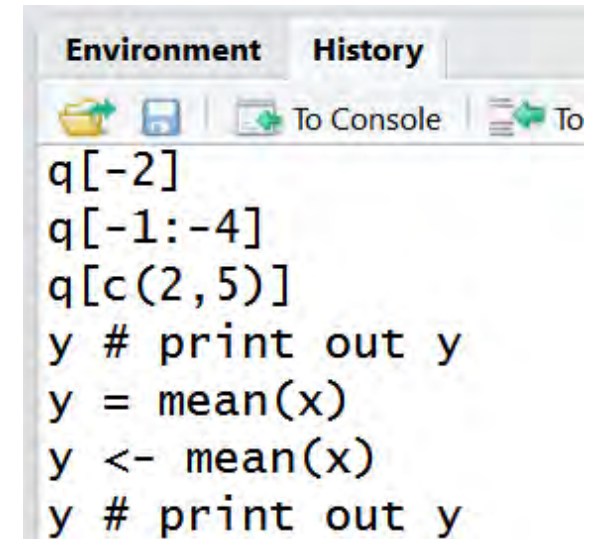
Comment

Try this

```
y # print out y
```

```
> y <- mean(x)
> y # print out y
[1] 2.333333
```

- Text after # is treated as comment
- They are not executed, but recorded in history



The screenshot shows the R Studio interface with the 'Environment' and 'History' panes. The 'History' pane is active, displaying a list of executed commands. The commands are: q[-2], q[-1:-4], q[c(2,5)], y # print out y, y = mean(x), y <- mean(x), and y # print out y. The commands are listed in a scrollable area with icons for saving, deleting, and other actions.

Internal Dataset

- R is preloaded with the internal datasets
- To view the whole list type `data()`

```
R data sets x
Data sets in package 'datasets':

AirPassengers      Monthly Airline Passenger Numbers 1949-1960
BJsales            Sales Data with Leading Indicator
BJsales.lead (BJsales) Sales Data with Leading Indicator
BOD                Biochemical Oxygen Demand
CO2                Carbon Dioxide Uptake in Grass Plants
ChickWeight        Weight versus age of chicks on different diets
DNase              Elisa assay of DNase
EuStockMarkets     Daily Closing Prices of Major European Stock Indices, 1991-1998
Formaldehyde       Determination of Formaldehyde
HairEyeColor       Hair and Eye Color of Statistics Students
Harman23.cor        Harman Example 2.3
Harman74.cor        Harman Example 7.4
Indometh            Pharmacokinetics of Indomethacin
InsectSprays       Effectiveness of Insect Sprays
JohnsonJohnson    Quarterly Earnings per Johnson & Johnson
```


Internal Dataset


- You can load the internal dataset by typing the name
- Try the following

`Nile`

`iris`

`mtcars`

- To plot a histogram of Nile d
`hist(Nile)`



Help

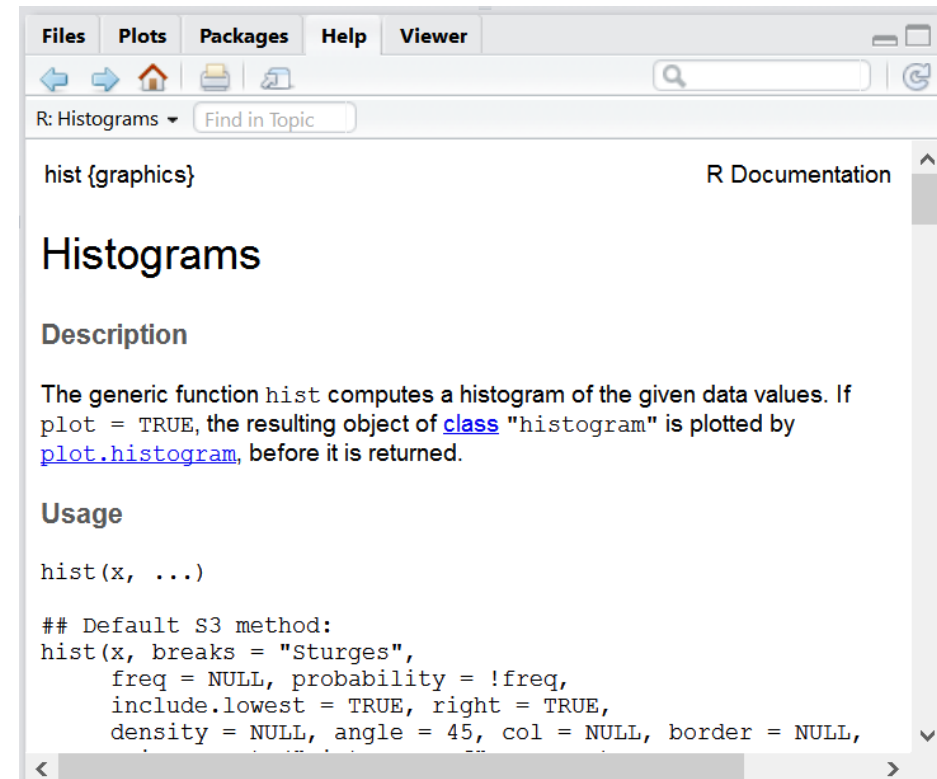
- A skilled programmer always look for documentation
- To access the function, data or package documentations, try

```
help( "hist" )
```

or

```
?hist
```

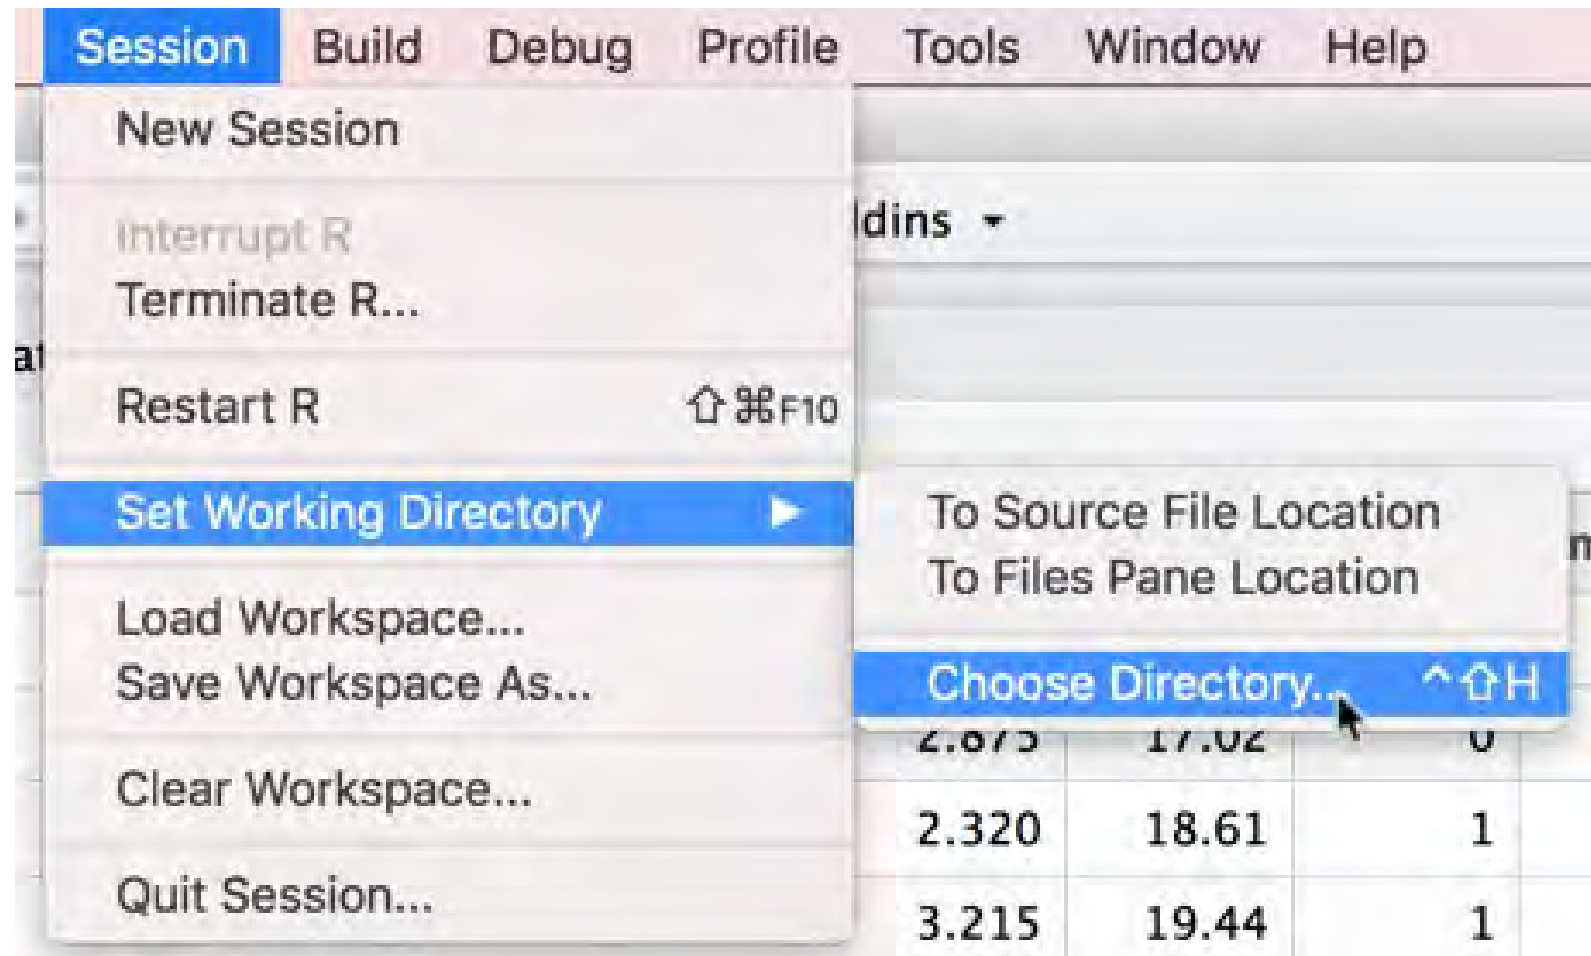
```
?Nile
```



Working directory

- Working directory is the current directory that the console prompt is in.
- Working directory is like Linux or Windows command line current directory.
- R users need to set working directory in order to use the relative path or using the file directory without path.

Change working directory



Read Data

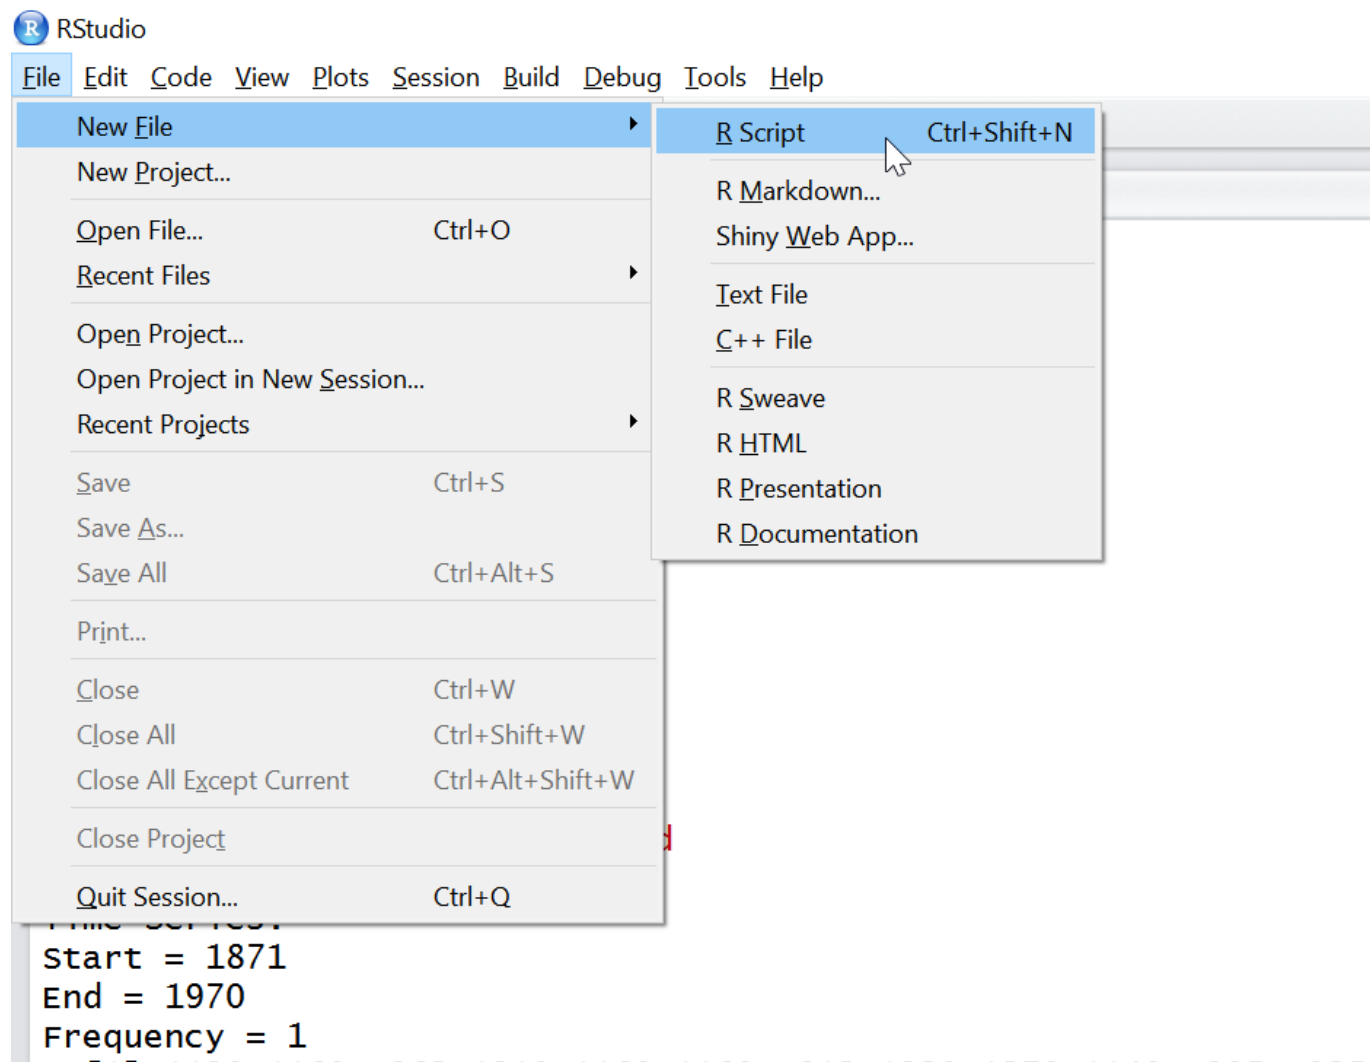
```
> flightData <- read.csv("flights.csv")  
> names(flightData)
```

| | | |
|------------------|-------------|------------------|
| [1] "X" | "year" | "month" |
| [4] "day" | "dep_time" | "sched_dep_time" |
| [7] "dep_delay" | "arr_time" | "sched_arr_time" |
| [10] "arr_delay" | "carrier" | "flight" |
| [13] "tailnum" | "origin" | "dest" |
| [16] "air_time" | "distance" | "hour" |
| [19] "minute" | "time_hour" | |

Introduction to function

- As in most programming languages, the heart of R programming consists of writing *functions*.
- A function is a group of instructions that takes inputs, uses them to compute other values, and returns a result.
- As a simple introduction, let's define a function named `oddcount()`, whose purpose is to count the odd numbers in a vector of integers.
- We will write a function in RStudio editor, save it to file and load it to the workspace.

1. Create a new R script file

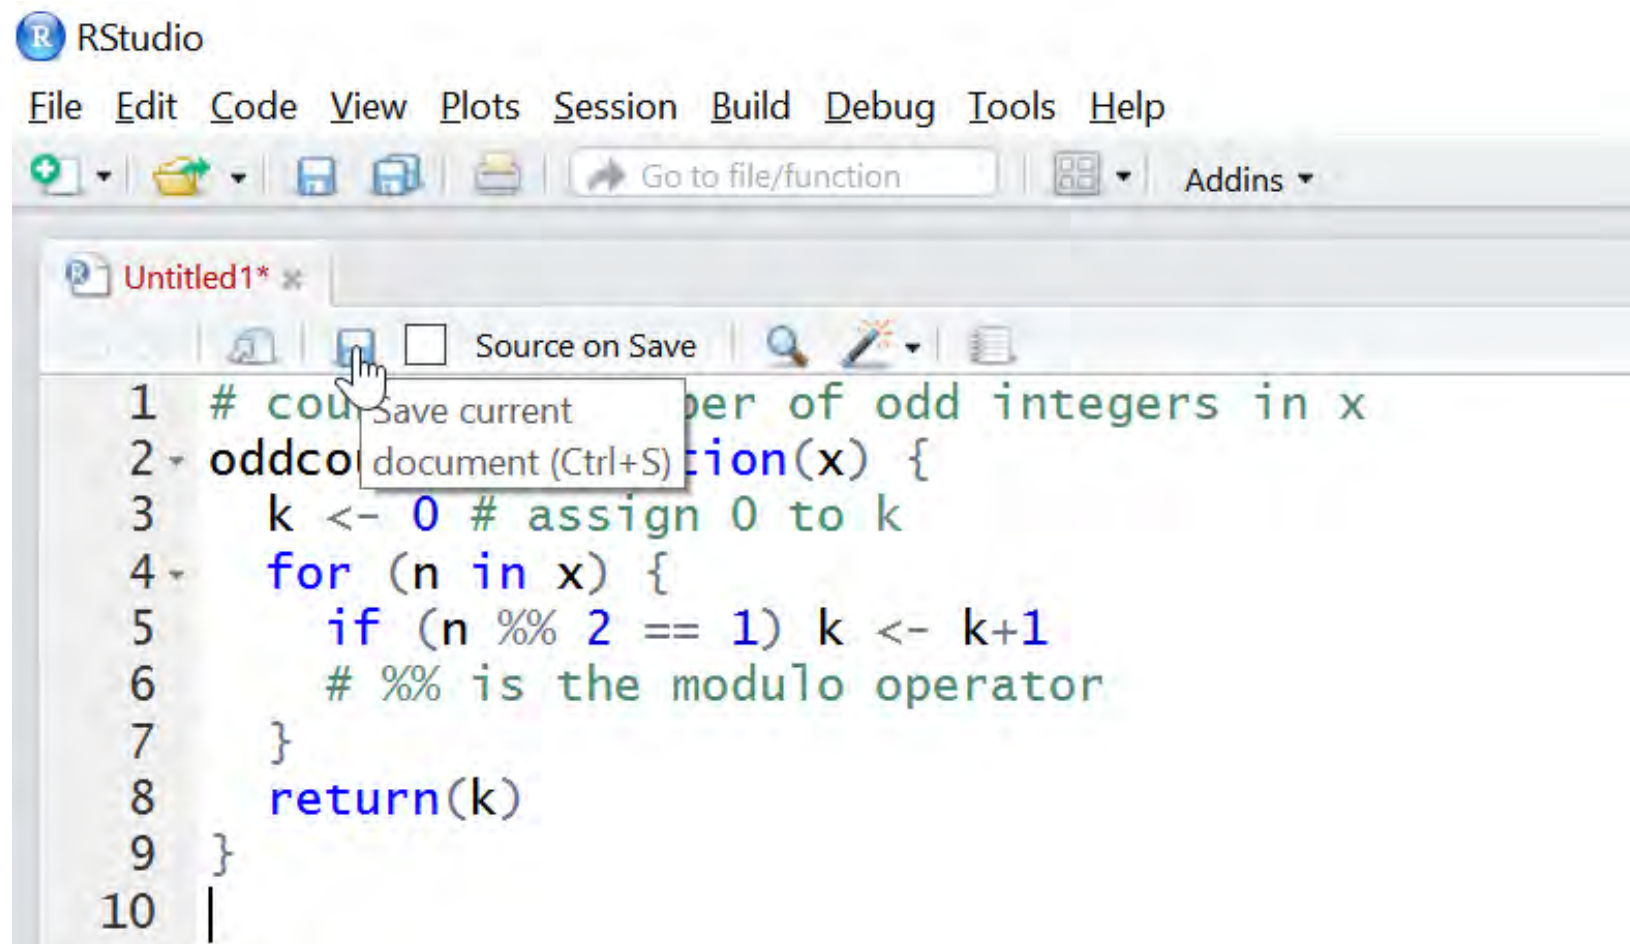


2. Write the oddcount function

counts the number of odd integers in x

```
oddcount <- function(x) {  
  k <- 0 # assign 0 to k  
  for (n in x) {  
    if (n %% 2 == 1) k <- k+1  
    # %% is the modulo operator  
  }  
  return(k)  
}
```

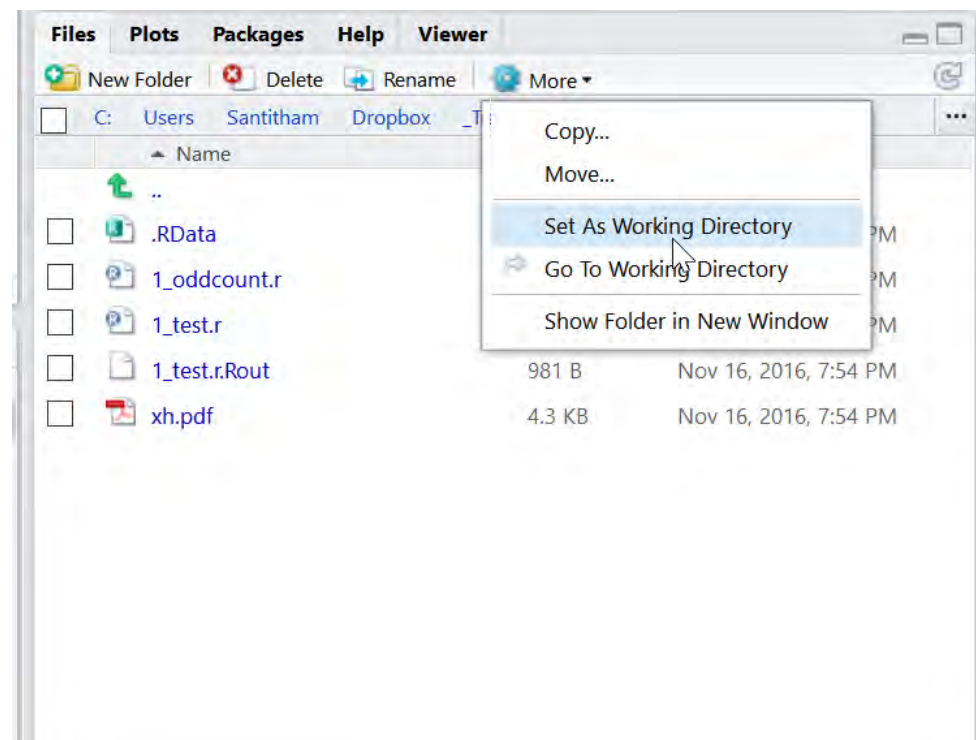
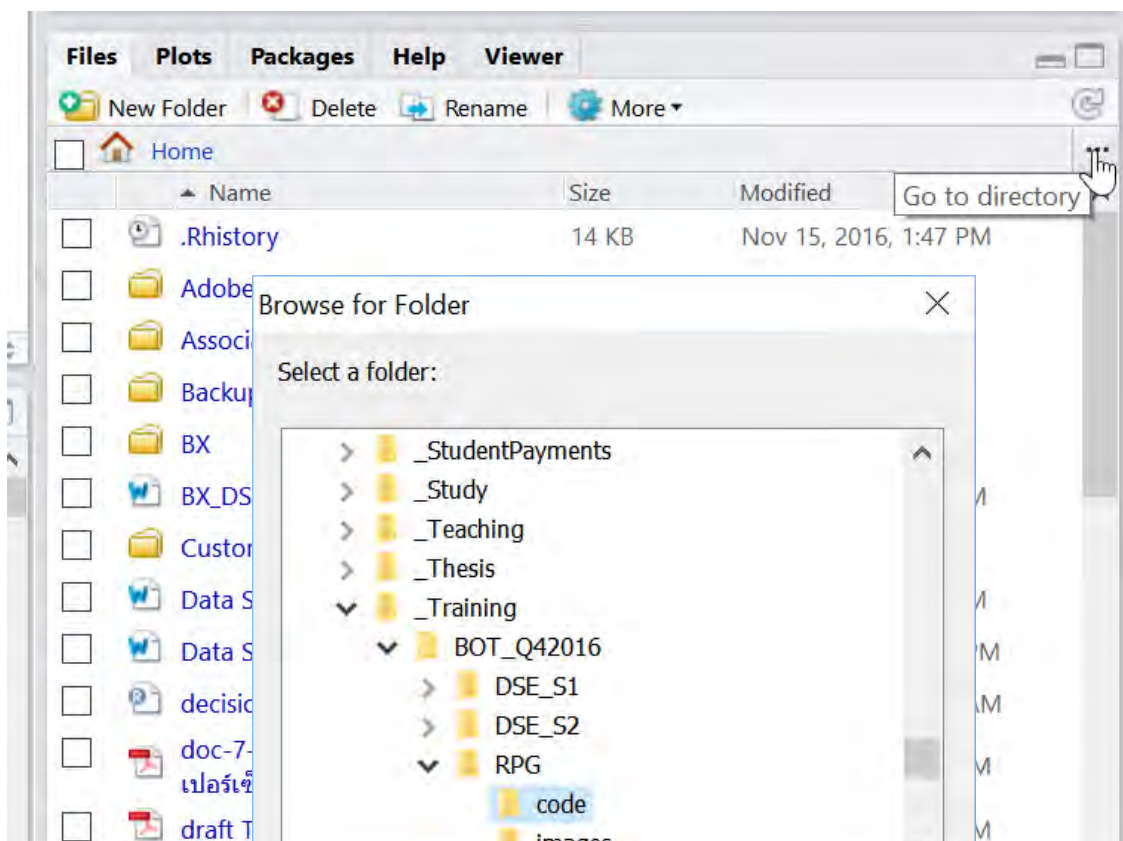

3. Save file with name “1_oddcount.r”



The screenshot shows the RStudio application window. The menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Tools, and Help. The toolbar contains icons for creating a new file, opening a file, saving a file, printing, and a search bar labeled 'Go to file/function'. Below the toolbar, a tab labeled 'Untitled1*' is active. The script editor shows R code for a function named 'oddcount'. A mouse cursor is hovering over the 'Save' icon in the toolbar, which has opened a tooltip that reads 'Save current document (Ctrl+S)'. The code in the editor is as follows:

```
1 # count the number of odd integers in x
2 oddcount <- function(x) {
3   k <- 0 # assign 0 to k
4   for (n in x) {
5     if (n %% 2 == 1) k <- k+1
6     # %% is the modulo operator
7   }
8   return(k)
9 }
10 |
```

4. Change working directory to where you save file



```
> setwd("C:/Users/Santitham/Dropbox/_Training/BOT_Q42016/RPG/code")
```

```
> |
```

5. Load the function to use

To load the function

```
source( "1_oddcount.r" )
```

To use the function

```
oddcount( c( 1, 2, 3, 5, 7, 10 ) )
```

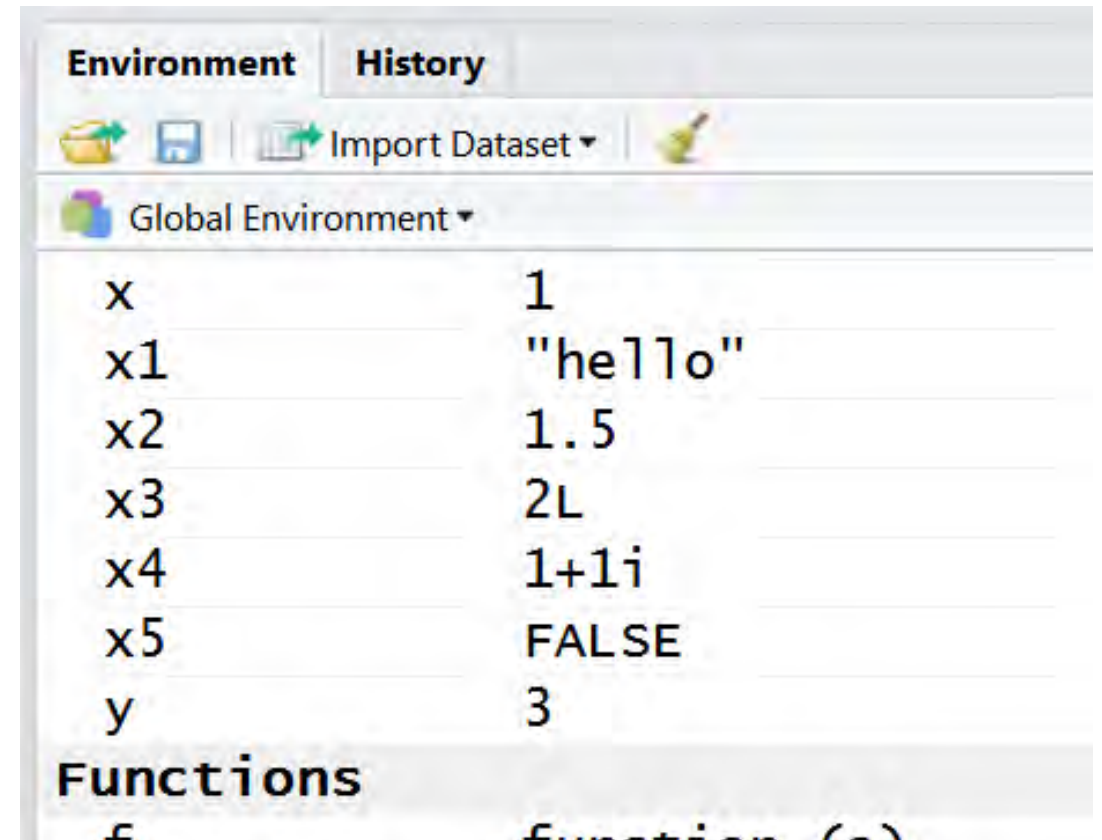
```
> oddcount(c(1,2,3,5,7,10))  
[1] 4
```





R objects – basic classes

- R has five basic classes of objects
 - character
 - numeric (real numbers)
 - integer
 - complex
 - logical (True/False)
- Special values
 - Inf. Infinity occurs when the number is divided by zero
 - NaN. Not-a-number indicates undefined value, e.g. 0/0. It can also be thought as a missing value.

Try this

```
x1 <- 'hello'
x2 <- 1.5
x3 <- as.integer(2)
x4 <- 1+1i
x5 <- FALSE
```



| Environment | | History |
|--|-------------|---|
|    Import Dataset ▾ | |  |
| Global Environment ▾ | | |
| x | 1 | |
| x1 | "hello" | |
| x2 | 1.5 | |
| x3 | 2L | |
| x4 | 1+1i | |
| x5 | FALSE | |
| y | 3 | |
| Functions | | |
| f | function () | |

Compound data structure

- Vector
- Matrices and Arrays
- Lists
- Data Frames
- Factors and Tables

Vectors

- Vector is a basic foundation of other data structure

```
x <- 1
```

- This generates one-element vector

```
x <- c(2, 4, 6)
```

- This generates three-elements vector

Adding or deleting vector elements

- Vectors are stored contiguously, thus you cannot insert or delete elements
- The size of a vector is determined at its creation, so if you wish to add or delete elements, you'll need to reassign the vector.

```
> x <- 1:5  
> x <- c(x[1:3],0,x[4:5])  
> x  
[1] 1 2 3 0 4 5  
>  
> length(x)  
[1] 6
```


Common vector operation

Arithmetic operator

- Adding/subtracting

```
x <- c(1, 2, 4)
```

```
x + c(5, 0, 1)
```

```
x - 1
```

```
x + c(2, 1)
```

- (Element-wise) multiplication

```
x * c(5, 0, 4)
```

- (Element-wise) division

```
x / c(5, 4, -1)
```

- (Element-wise) modulus

```
x %% c(5, 4, -1)
```

NA: missing data

```
> x <- c(1:5,NA,6:10)
```

```
> x
```

```
[1] 1 2 3 4 5 NA 6 7 8 9 10
```

```
> mean(x)
```

```
[1] NA
```

```
> mean(x, na.rm = TRUE)
```

```
[1] 5.5
```

```
> mean(na.omit(x))
```

```
[1] 5.5
```

NULL: nothing, empty set

```
> x1 <- c(1:5,NULL,6:10)
> x1
[1] 1 2 3 4 5 6 7 8 9 10
```

- NULL can be used as an initial empty variable

```
> z <- NULL
> for (i in 1:10) if (i%%2==0) z <- c(z,i)
> z
[1] 2 4 6 8 10
```

Filtering

- Filtering allows us to extract a vector's elements that satisfy certain conditions.
- This is because the index vector become a logical vector

```
> x <- 1:10  
> x[x > 5]  
[1] 6 7 8 9 10
```

```
> x > 5  
[1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
```

- We can assign value to the filtered position

```
> x[x<4] <- 0  
> x  
[1] 0 0 0 4 5 6 7 8 9 10
```

Selection function which()

- In some cases, though, we may just want to find the positions within a vector at which the condition occurs.
- We can do this using `which()`, as follows:

```
> x <- 1:10
> x
[1] 1 2 3 4 5 6 7 8 9 10
> which(x*x > 50)
[1] 8 9 10
```

Matrix Creating matrix

```
> m <- matrix(1:4,nrow=2,ncol=2)
```

```
> m
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 1 | 3 |
| [2,] | 2 | 4 |

```
> y <- matrix(1:6,nrow=2)
```

```
> y
```

| | [,1] | [,2] | [,3] |
|------|------|------|------|
| [1,] | 1 | 3 | 5 |
| [2,] | 2 | 4 | 6 |

```
> y <- matrix(1:6,nrow=2,byrow = T)
```

```
> y
```

| | [,1] | [,2] | [,3] |
|------|------|------|------|
| [1,] | 1 | 2 | 3 |
| [2,] | 4 | 5 | 6 |

Matrix operation

Matrix Multiplication by Scalar

```
> m
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 1 | 3 |
| [2,] | 2 | 4 |

```
> m %*% m
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 7 | 15 |
| [2,] | 10 | 22 |

Matrix Multiplication

```
> 3*m
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 3 | 9 |
| [2,] | 6 | 12 |

```
> m + m
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 2 | 6 |
| [2,] | 4 | 8 |

Matrix Addition

Matrix indexing

```
> z <- matrix(1:16, nrow = 4)
```

```
> z
```

| | [,1] | [,2] | [,3] | [,4] |
|------|------|------|------|------|
| [1,] | 1 | 5 | 9 | 13 |
| [2,] | 2 | 6 | 10 | 14 |
| [3,] | 3 | 7 | 11 | 15 |
| [4,] | 4 | 8 | 12 | 16 |

```
> z[,2:3]
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 5 | 9 |
| [2,] | 6 | 10 |
| [3,] | 7 | 11 |
| [4,] | 8 | 12 |

Selecting only 2nd and
3rd columns

Adding or deleting matrix rows and columns

- Just like vector, matrices are fixed length and dimension
- However, they can be reassigned
- `cbind()` bind the column together

```
> z
      [,1] [,2]
[1,]    1    4
[2,]    2    5
[3,]    3    6
> cbind(z, rep(1,3))
      [,1] [,2] [,3]
[1,]    1    4    1
[2,]    2    5    1
[3,]    3    6    1
```

Object coercion

If more than one class are introduced, the vector will be coerced.

Implicit

```
> y <- c(1.7, "a")
```

```
## character
```

```
> y <- c(TRUE, 2)
```

```
## numeric
```

```
> y <- c("a", TRUE)
```

```
## character
```

Explicit

```
> x <- 0:6
```

```
> class(x)
```

```
[1] "integer"
```

```
> as.numeric(x)
```

```
[1] 0 1 2 3 4 5 6
```

```
> as.logical(x)
```

```
[1] FALSE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
> as.character(x)
```

```
[1] "0" "1" "2" "3" "4" "5" "6"
```

List

- In contrast to a vector, in which all elements must be of the same mode, R's list structure can combine objects of different types.
- Ordinary vectors are termed *atomic* vectors, since their components cannot be broken down into smaller components.
- In contrast, lists are referred to as *recursive* vectors.

Creating a list

- Let us consider an employee database.
- For each employee, we wish to store the name, salary, and a Boolean indicating union membership.

```
> j <- list(name="Joe", salary=55000, union=T)
```

```
> j
```

```
$name
```

```
[1] "Joe"
```

```
$salary
```

```
[1] 55000
```

```
$union
```

```
[1] TRUE
```

List indexing

- You can access a list component in several different ways:

```
> j$salary  
[1] 55000  
> j[["salary"]]  
[1] 55000  
> j[[2]]  
[1] 55000
```

Adding list elements

Element can be added directly



```
> j$location <- 'Bangkok'
> j
$name
[1] "Joe"

$salary
[1] 55000

$union
[1] TRUE

$location
[1] "Bangkok"
```

Deleting list elements

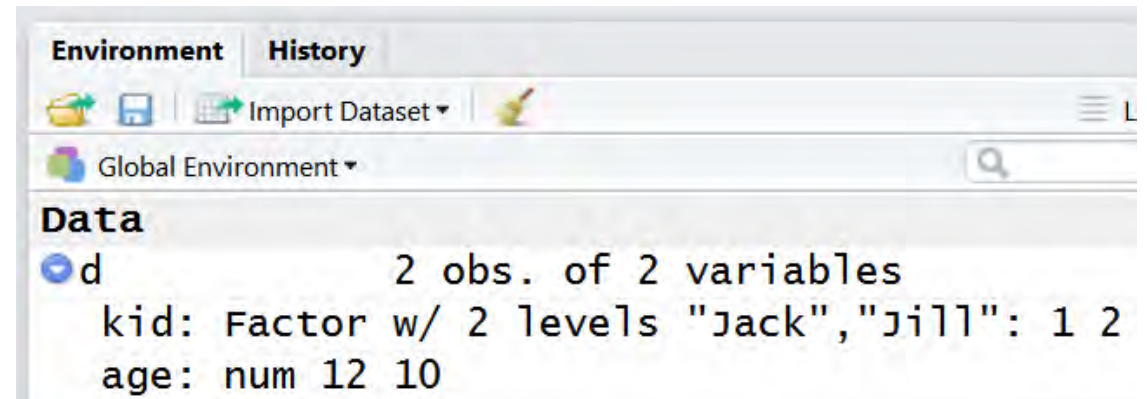
```
> j$location <- NULL  
> j  
$name  
[1] "Joe"  
  
$salary  
[1] 55000  
  
$union  
[1] TRUE
```

Data Frames

- On an intuitive level, a *data frame* is like a matrix, with a two-dimensional rows and columns structure.
- On a technical level, a data frame is a list, with the components of that list being equal-length vectors.

```
> kid <- c("Jack","Jill")
> age <- c(12,10)
> d <- data.frame(kid,age)
> d
```

| | kid | age |
|---|------|-----|
| 1 | Jack | 12 |
| 2 | Jill | 10 |



Accessing data frame

```
> d[[1]]  
[1] Jack Jill  
Levels: Jack Jill  
> d$kid  
[1] Jack Jill  
Levels: Jack Jill  
> d[,1]  
[1] Jack Jill  
Levels: Jack Jill  
,
```

Object internal structure

```
> str(d)
'data.frame':   2 obs. of  2 variables:
 $ kid: Factor w/ 2 levels "Jack","Jill": 1 2
 $ age: num  12 10
```

Example: mtcars

- You can view the top rows of the data frame by using `head()`

```
> head(mtcars)
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| Valiant | 18.1 | 6 | 225 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

Adding column to data frame

```
> mtcars$ratio = mtcars$hp / mtcars$cyl
> head(mtcars)
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| Valiant | 18.1 | 6 | 225 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

| | ratio |
|-------------------|----------|
| Mazda RX4 | 18.33333 |
| Mazda RX4 Wag | 18.33333 |
| Datsun 710 | 23.25000 |
| Hornet 4 Drive | 18.33333 |
| Hornet Sportabout | 21.87500 |
| Valiant | 17.50000 |

Delete column from data frame

```
> mtcars$ratio <- NULL  
> head(mtcars)
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| valiant | 18.1 | 6 | 225 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

Adding row to data frame

```
> tail(rbind(mtcars,list(99,10,200,999,9,9,99,1,1,9,9)))
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|----------------|------|-----|-------|-----|------|-------|------|----|----|------|------|
| Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 | 1.513 | 16.9 | 1 | 1 | 5 | 2 |
| Ford Pantera L | 15.8 | 8 | 351.0 | 264 | 4.22 | 3.170 | 14.5 | 0 | 1 | 5 | 4 |
| Ferrari Dino | 19.7 | 6 | 145.0 | 175 | 3.62 | 2.770 | 15.5 | 0 | 1 | 5 | 6 |
| Maserati Bora | 15.0 | 8 | 301.0 | 335 | 3.54 | 3.570 | 14.6 | 0 | 1 | 5 | 8 |
| Volvo 142E | 21.4 | 4 | 121.0 | 109 | 4.11 | 2.780 | 18.6 | 1 | 1 | 4 | 2 |
| 33 | 99.0 | 10 | 200.0 | 999 | 9.00 | 9.000 | 99.0 | 1 | 1 | 9 | 9 |

`rbind()` produces a new data frame with additional row
`tail()` function can be used to view the end of data frame

Factors

- Factors form the basis for many of R's powerful operations, including many of those performed on tabular data.
- The motivation for factors comes from the notion of *nominal*, or *categorical*, variables in statistics.
- These values are nonnumerical in nature, corresponding to categories such as RENT, MORTGAGE, OWN, although they may be coded using numbers.
- Statistical model relies on factors (and levels) to indicate group of data to be analyzed

Factors and levels

- An R *factor* might be viewed simply as a vector with a bit more information added (though, as seen below, it's different from this internally).
- That extra information consists of a record of the distinct values in that vector, called *levels*. Here's an example:

```
> design <- c('H','M','L','H','M','L')
> design
[1] "H" "M" "L" "H" "M" "L"
> design_factor <- factor(design)
> design_factor
[1] H M L H M L
Levels: H L M
```


Data Preparation

Section 2

Two packages to help you
work with the structure of data.



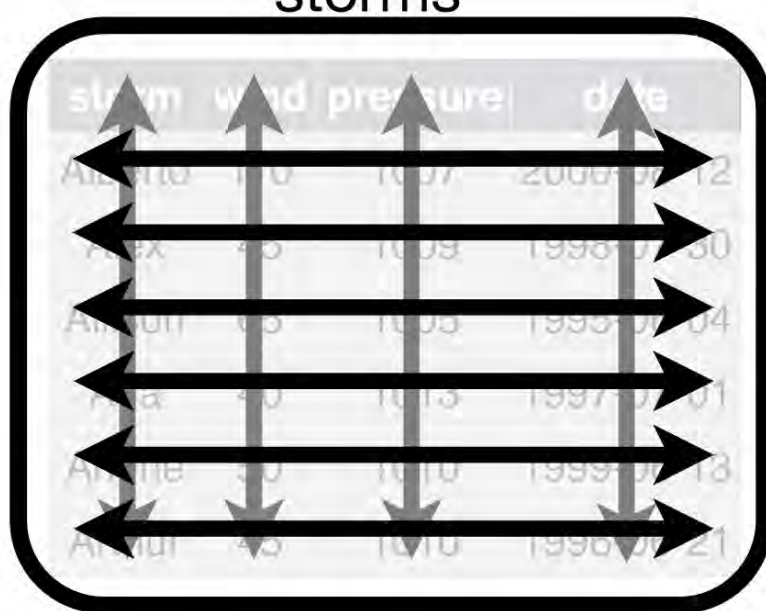
tidyr



dplyr

Tidy data

storms



| storm | wind | pressure | date |
|--------|------|----------|------------|
| Albino | 100 | 1007 | 2000-07-12 |
| Felix | 45 | 1009 | 1993-07-30 |
| Amelia | 65 | 1005 | 1993-07-04 |
| Fa | 40 | 1013 | 1997-07-01 |
| Alma | 50 | 1010 | 1999-07-18 |
| Amelia | 45 | 1010 | 1996-07-21 |

- 1** Each **variable** is saved in its own **column**.
- 2** Each **observation** is saved in its own **row**.
- 3** Each "type" of observation stored in a **single table** (here, storms).

Load data

Data

| | |
|-----------|--------------------------|
| cases | 3 obs. of 4 variables |
| pollution | 6 obs. of 3 variables |
| storms | 6 obs. of 4 variables |
| tb | 3800 obs. of 6 variables |

```
load("2-DataPrep.RData")
```

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |

cases

| Country | 2011 | 2012 | 2013 |
|---------|-------|-------|-------|
| FR | 7000 | 6900 | 7000 |
| DE | 5800 | 6000 | 6200 |
| US | 15000 | 14000 | 13000 |

pollution

| city | particle size | amount ($\mu\text{g}/\text{m}^3$) |
|----------|---------------|--|
| New York | large | 23 |
| New York | small | 14 |
| London | large | 22 |
| London | small | 16 |
| Beijing | large | 121 |
| Beijing | small | 56 |



Tidy data

Data Wrangling with dplyr and tidyr

Cheat Sheet



Syntax - Helpful conventions for wrangling

dplyr::tbl_df(iris)

Converts data to tbl class. tbl's are easier to examine than data frames. R displays only the data that fits onscreen:

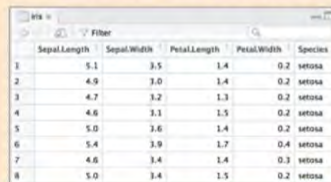
```
Source: local data frame [150 x 5]
  Sepal.Length Sepal.Width Petal.Length
1           5.1           3.5           1.4
2           4.9           3.0           1.4
3           4.7           3.2           1.3
4           4.6           3.1           1.5
5           5.0           3.6           1.4
...
Variables not shown: Petal.Width (dbl),
Species (fctr)
```

dplyr::glimpse(iris)

Information dense summary of tbl data.

utils::View(iris)

View data set in spreadsheet-like display (note capital V).



dplyr::%>%

Passes object on left hand side as first argument (or argument) of function on righthand side.

x %>% f(y) is the same as **f(x, y)**
y %>% f(x, .., z) is the same as **f(x, y, z)**

"Piping" with %>% makes code more readable, e.g.

```
iris %>%
  group_by(Species) %>%
  summarise(avg = mean(Sepal.Width)) %>%
  arrange(avg)
```

RStudio® is a trademark of RStudio, Inc. • All rights reserved • info@rstudio.com • 844-448-1212 • rstudio.com


Logic in R - ?Comparison, ?base::Logic

| | | | |
|----|--------------------------|------------------------|-------------------|
| < | Less than | != | Not equal to |
| > | Greater than | %in% | Group membership |
| == | Equal to | is.na | is NA |
| <= | Less than or equal to | !is.na | Is not NA |
| >= | Greater than or equal to | 4, , !, xor, any, all | Boolean operators |


devtools::install_github("rstudio/EDAWR") for updates.

Learn more with [browseVignettes\(package = c\("dplyr", "tidyr"\)\)](#) • dplyr 0.4.0 • tidyr 0.2.0 • Updated: 1/15

Tidy Data - A foundation for wrangling in R



Reshaping Data - Change the layout of a data set



Subset Observations (Rows)



dplyr::filter(iris, Sepal.Length > 7)

Extract rows that meet logical criteria.

dplyr::distinct(iris)

Remove duplicate rows.

dplyr::sample_frac(iris, 0.5, replace = TRUE)

Randomly select fraction of rows.

dplyr::sample_n(iris, 10, replace = TRUE)

Randomly select n rows.

dplyr::slice(iris, 10:15)

Select rows by position.

dplyr::top_n(storms, 2, date)

Select and order top n entries (by group if grouped data).

Subset Variables (Columns)



dplyr::select(iris, Sepal.Width, Petal.Length, Species)

Select columns by name or helper function.

Helper functions for select - ?select

select(iris, contains(" "))
Select columns whose name contains a character string.

select(iris, ends_with("Length"))
Select columns whose name ends with a character string.

select(iris, everything())
Select every column.

select(iris, matches("t"))
Select columns whose name matches a regular expression.

select(iris, num_range("x", 1:5))
Select columns named x1, x2, x3, x4, x5.

select(iris, one_of(c("Species", "Genus")))
Select columns whose names are in a group of names.

select(iris, starts_with("Sepal"))
Select columns whose name starts with a character string.

select(iris, Sepal.Length:Petal.Width)
Select all columns between Sepal.Length and Petal.Width (inclusive).

select(iris, -Species)
Select all columns except Species.

Data Preparation Tasks

- Reshaping table: gather, spread
- Reshaping column: unite, separate
- Subsetting rows: filter, distinct, sample_n, top_n
- Subsetting columns: select
- Summarise data: summarise
- Group data: group_by
- Make new variables: mutate
- Integration: left_join, union, bind_cols

tidyr

Gather columns

| ## | country | 2011 | 2012 | 2013 |
|------|---------|-------|-------|-------|
| ## 1 | FR | 7000 | 6900 | 7000 |
| ## 2 | DE | 5800 | 6000 | 6200 |
| ## 3 | US | 15000 | 14000 | 13000 |



| ## | country | year | n |
|------|---------|------|-------|
| ## 1 | FR | 2011 | 7000 |
| ## 2 | DE | 2011 | 5800 |
| ## 3 | US | 2011 | 15000 |
| ## 4 | FR | 2012 | 6900 |
| ## 5 | DE | 2012 | 6000 |
| ## 6 | US | 2012 | 14000 |
| ## 7 | FR | 2013 | 7000 |
| ## 8 | DE | 2013 | 6200 |
| ## 9 | US | 2013 | 13000 |

```
gather(cases, "year", "n", 2:4)
```

tidyr

Spread columns

| ## | | city | size | amount |
|------|----------|-------|------|--------|
| ## 1 | New York | large | 23 | |
| ## 2 | New York | small | 14 | |
| ## 3 | London | large | 22 | |
| ## 4 | London | small | 16 | |
| ## 5 | Beijing | large | 121 | |
| ## 6 | Beijing | small | 56 | |



| ## | | city | large | small |
|------|----------|------|-------|-------|
| ## 1 | Beijing | 121 | 56 | |
| ## 2 | London | 22 | 16 | |
| ## 3 | New York | 23 | 14 | |

```
spread(pollution, size, amount)
```


separate()

Separate splits a column by a character string separator.

```
separate(storms, date, c("year", "month", "day"), sep = "-")
```

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



storms2

| storm | wind | pressure | year | month | day |
|---------|------|----------|------|-------|-----|
| Alberto | 110 | 1007 | 2000 | 08 | 12 |
| Alex | 45 | 1009 | 1998 | 07 | 30 |
| Allison | 65 | 1005 | 1995 | 06 | 04 |
| Ana | 40 | 1013 | 1997 | 07 | 1 |
| Arlene | 50 | 1010 | 1999 | 06 | 13 |
| Arthur | 45 | 1010 | 1996 | 06 | 21 |

unite()

Unite unites columns into a single column.

```
unite(storms2, "date", year, month, day, sep = "-")
```

storms2

| storm | wind | pressure | year | month | day |
|---------|------|----------|------|-------|-----|
| Alberto | 110 | 1007 | 2000 | 08 | 12 |
| Alex | 45 | 1009 | 1998 | 07 | 30 |
| Allison | 65 | 1005 | 1995 | 06 | 04 |
| Ana | 40 | 1013 | 1997 | 07 | 1 |
| Arlene | 50 | 1010 | 1999 | 06 | 13 |
| Arthur | 45 | 1010 | 1996 | 06 | 21 |



storms

| storm | wind | pressure | date | | |
|---------|------|----------|------------|--|--|
| Alberto | 110 | 1007 | 2000-08-12 | | |
| Alex | 45 | 1009 | 1998-07-30 | | |
| Allison | 65 | 1005 | 1995-06-04 | | |
| Ana | 40 | 1013 | 1997-07-01 | | |
| Arlene | 50 | 1010 | 1999-06-13 | | |
| Arthur | 45 | 1010 | 1996-06-21 | | |

Ways to access information

- 1** **Extract** existing variables. **select()**
- 2** **Extract** existing observations. **filter()**
- 3** **Derive** new variables
(from existing variables) **mutate()**
- 4** **Change** the unit of analysis **summarise()**

select()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | pressure |
|---------|----------|
| Alberto | 1007 |
| Alex | 1009 |
| Allison | 1005 |
| Ana | 1013 |
| Arlene | 1010 |
| Arthur | 1010 |

```
select(storms, storm, pressure)
```

select()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| wind | pressure | date |
|------|----------|------------|
| 110 | 1007 | 2000-08-12 |
| 45 | 1009 | 1998-07-30 |
| 65 | 1005 | 1995-06-04 |
| 40 | 1013 | 1997-07-01 |
| 50 | 1010 | 1999-06-13 |
| 45 | 1010 | 1996-06-21 |

```
select(storms, -storm)
```

```
# see ?select for more
```

select()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| wind | pressure | date |
|------|----------|------------|
| 110 | 1007 | 2000-08-12 |
| 45 | 1009 | 1998-07-30 |
| 65 | 1005 | 1995-06-04 |
| 40 | 1013 | 1997-07-01 |
| 50 | 1010 | 1999-06-13 |
| 45 | 1010 | 1996-06-21 |

```
select(storms, wind:date)
```

```
# see ?select for more
```


Useful select functions

* Blue functions come in dplyr

| | |
|---------------|--|
| - | Select everything but |
| : | Select range |
| contains() | Select columns whose name contains a character string |
| ends_with() | Select columns whose name ends with a string |
| everything() | Select every column |
| matches() | Select columns whose name matches a regular expression |
| num_range() | Select columns named x1, x2, x3, x4, x5 |
| one_of() | Select columns whose names are in a group of names |
| starts_with() | Select columns whose name starts with a character string |

filter()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Allison | 65 | 1005 | 1995-06-04 |
| Arlene | 50 | 1010 | 1999-06-13 |

```
filter(storms, wind >= 50)
```


filter()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Allison | 65 | 1005 | 1995-06-04 |

```
filter(storms, wind >= 50,  
       storm %in% c("Alberto", "Alex", "Allison"))
```

logical tests in R

?Comparison

| | |
|--------|--------------------------|
| < | Less than |
| > | Greater than |
| == | Equal to |
| <= | Less than or equal to |
| >= | Greater than or equal to |
| != | Not equal to |
| %in% | Group membership |
| is.na | Is NA |
| !is.na | Is not NA |

?base::Logic

| | |
|-----|-------------|
| & | boolean and |
| | boolean or |
| xor | exactly or |
| ! | not |
| any | any true |
| all | all true |

mutate()

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | wind | pressure | date | ratio | inverse |
|---------|------|----------|------------|-------|---------|
| Alberto | 110 | 1007 | 2000-08-12 | 9.15 | 0.11 |
| Alex | 45 | 1009 | 1998-07-30 | 22.42 | 0.04 |
| Allison | 65 | 1005 | 1995-06-04 | 15.46 | 0.06 |
| Ana | 40 | 1013 | 1997-07-01 | 25.32 | 0.04 |
| Arlene | 50 | 1010 | 1999-06-13 | 20.20 | 0.05 |
| Arthur | 45 | 1010 | 1996-06-21 | 22.44 | 0.04 |

```
mutate(storms, ratio = pressure / wind, inverse = ratio^-1)
```


Useful mutate functions

* All take a vector of values and return a vector of values

** Blue functions come in dplyr

| | | |
|---|-----------------------------------|--|
| pmin(), pmax() | Element-wise min and max | |
| cummin(), cummax() | Cumulative min and max | |
| cumsum(), cumprod() | Cumulative sum and product | |
| between() | Are values between a and b? | |
| cume_dist() | Cumulative distribution of values | |
| cumall(), cumany() | Cumulative all and any | |
| cummean() | Cumulative mean | |
| lead(), lag() | Copy with values one position | |
| ntile() | Bin vector into n buckets | |
| dense_rank(), min_rank(), percent_rank(), row_number() | Various ranking methods | |

"Window" functions

* All take a vector of values and return a vector of values

pmin(), pmax()
cummin(), cummax()
cumsum(), cumprod()

between()

cume_dist()

cumall(), cumany()

cummean()

lead(), lag()

ntile()

dense_rank(), min_rank(),
percent_rank(), row_number()

1
2
3
4
5
6

cumsum()

1
3
6
10
15
21

summarise()

| city | particle size | amount ($\mu\text{g}/\text{m}^3$) |
|----------|---------------|--|
| New York | large | 23 |
| New York | small | 14 |
| London | large | 22 |
| London | small | 16 |
| Beijing | large | 121 |
| Beijing | small | 56 |



| median | variance |
|--------|----------|
| 22.5 | 1731.6 |

```
pollution %>% summarise(median = median(amount), variance = var(amount))
```


summarise()

| city | particle size | amount ($\mu\text{g}/\text{m}^3$) |
|----------|---------------|--|
| New York | large | 23 |
| New York | small | 14 |
| London | large | 22 |
| London | small | 16 |
| Beijing | large | 121 |
| Beijing | small | 56 |



| mean | sum | n |
|------|-----|---|
| 42 | 252 | 6 |

```
pollution %>% summarise(mean = mean(amount), sum = sum(amount), n = n())
```

Useful summary functions

* All take a vector of values and return a single value

** Blue functions come in dplyr

| | | |
|--------------|---|--|
| min(), max() | Minimum and maximum values | |
| mean() | Mean value | |
| median() | Median value | |
| sum() | Sum of values | |
| var, sd() | Variance and standard deviation of a vector | |
| first() | First value in a vector | |
| last() | Last value in a vector | |
| nth() | Nth value in a vector | |
| n() | The number of values in a vector | |
| n_distinct() | The number of distinct values in a vector | |

"Summary" functions

* All take a vector of values and return a single value

min(), max()

mean()

median()

sum()

var, sd()

first()

last()

nth()

n()

n_distinct()

1

2

3

4

5

6

sum()

21

arrange()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | wind | pressure | date |
|---------|------|----------|------------|
| Ana | 40 | 1013 | 1997-07-01 |
| Alex | 45 | 1009 | 1998-07-30 |
| Arthur | 45 | 1010 | 1996-06-21 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Allison | 65 | 1005 | 1995-06-04 |
| Alberto | 110 | 1007 | 2000-08-12 |

arrange(storms, **wind**)

arrange()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | wind | pressure | date |
|---------|------|----------|------------|
| Ana | 40 | 1013 | 1997-07-01 |
| Alex | 45 | 1009 | 1998-07-30 |
| Arthur | 45 | 1010 | 1996-06-21 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Allison | 65 | 1005 | 1995-06-04 |
| Alberto | 110 | 1007 | 2000-08-12 |

arrange(storms, wind)

arrange()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |




| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Allison | 65 | 1005 | 1995-06-04 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |
| Alex | 45 | 1009 | 1998-07-30 |
| Ana | 40 | 1013 | 1997-07-01 |

`arrange(storms, desc(wind))`

The pipe operator $\%>\%$

```
library(dplyr)  
select(tb, child:elderly)  
tb %>% select(child:elderly)
```

These do the
same thing
Try it!



select()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | pressure |
|---------|----------|
| Alberto | 1007 |
| Alex | 1009 |
| Allison | 1005 |
| Ana | 1013 |
| Arlene | 1010 |
| Arthur | 1010 |

select(storms, storm, pressure)

select()

storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | pressure |
|---------|----------|
| Alberto | 1007 |
| Alex | 1009 |
| Allison | 1005 |
| Ana | 1013 |
| Arlene | 1010 |
| Arthur | 1010 |

```
storms %>% select(storm, pressure)
```


storms

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | pressure |
|---------|----------|
| Alberto | 1007 |
| Allison | 1005 |
| Arlene | 1010 |

```
storms %>%
```

```
  filter(wind >= 50) %>%
```

```
  select(storm, pressure)
```


mutate()

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



storms %>%

mutate(ratio = pressure / wind) %>%

select(storm, ratio)

mutate()

| storm | wind | pressure | date |
|---------|------|----------|------------|
| Alberto | 110 | 1007 | 2000-08-12 |
| Alex | 45 | 1009 | 1998-07-30 |
| Allison | 65 | 1005 | 1995-06-04 |
| Ana | 40 | 1013 | 1997-07-01 |
| Arlene | 50 | 1010 | 1999-06-13 |
| Arthur | 45 | 1010 | 1996-06-21 |



| storm | ratio |
|---------|-------|
| Alberto | 9.15 |
| Alex | 22.42 |
| Allison | 15.46 |
| Ana | 25.32 |
| Arlene | 20.20 |
| Arthur | 22.44 |

storms %>%


```
mutate(ratio = pressure / wind) %>%  
select(storm, ratio)
```

dplyr::**bind_cols()**

| y | | | z | | | | | | |
|----|----|---|----|----|---|----|----|----|----|
| x1 | x2 | | x1 | x2 | | x1 | x2 | x1 | x2 |
| A | 1 | + | B | 2 | = | A | 1 | B | 2 |
| B | 2 | | C | 3 | | B | 2 | C | 3 |
| C | 3 | | D | 4 | | C | 3 | D | 4 |

`bind_cols(y, z)`

dplyr::**bind_rows()**



| city | particle size | amount ($\mu\text{g}/\text{m}^3$) |
|----------|---------------|--|
| New York | large | 23 |
| New York | small | 14 |



| mean | sum | n |
|------|-----|---|
| 18.5 | 37 | 2 |

| | | |
|--------|-------|----|
| London | large | 22 |
| London | small | 16 |



| | | |
|------|----|---|
| 19.0 | 38 | 2 |
|------|----|---|

| | | |
|---------|-------|-----|
| Beijing | large | 121 |
| Beijing | small | 56 |



| | | |
|------|-----|---|
| 88.5 | 177 | 2 |
|------|-----|---|

`group_by() + summarise()`

| city | size | amount |
|----------|-------|--------|
| New York | large | 23 |
| New York | small | 14 |
| London | large | 22 |
| London | small | 16 |
| Beijing | large | 121 |
| Beijing | small | 56 |



| city | size | amount |
|----------|-------|--------|
| New York | large | 23 |
| New York | small | 14 |
| London | large | 22 |
| London | small | 16 |
| Beijing | large | 121 |
| Beijing | small | 56 |



| city | mean |
|----------|------|
| New York | 18.5 |
| London | 19.0 |
| Beijing | 88.5 |

```
pollution %>% group_by(city) %>% summarise(mean = mean(amount))
```

dplyr::**left_join()**

| songs | | | artists | | | | | |
|---------------------|-------|---|---------|--------|---|---------------------|-------|--------|
| song | name | | name | plays | | song | name | plays |
| Across the Universe | John | + | George | sitar | = | Across the Universe | John | guitar |
| Come Together | John | | John | guitar | | Come Together | John | guitar |
| Hello, Goodbye | Paul | | Paul | bass | | Hello, Goodbye | Paul | bass |
| Peggy Sue | Buddy | | Ringo | drums | | Peggy Sue | Buddy | <NA> |

```
left_join(songs, artists, by = "name")
```

Lab

- Load two data frames “flights.csv” and “airlines.csv”
 - Hint: `read.csv`
- Select two columns (carrier, dep_delay)
- Filter NA out of dep_delay
- Use carrier as a group and calculate mean of dep_delay
- Sort the worst carrier by departure delay (dep_delay)
- Join airlines data to flight data using ‘carrier’ as a key
- Show the result

Expected Results

A tibble: 16 x 3

| | carrier | mean_delay | name |
|----|---------|------------|-----------------------------|
| | <chr> | <dbl> | <chr> |
| 1 | F9 | 20.215543 | Frontier Airlines Inc. |
| 2 | EV | 19.955390 | ExpressJet Airlines Inc. |
| 3 | YV | 18.996330 | Mesa Airlines Inc. |
| 4 | FL | 18.726075 | AirTran Airways Corporation |
| 5 | WN | 17.711744 | Southwest Airlines Co. |
| 6 | 9E | 16.725769 | Endeavor Air Inc. |
| 7 | B6 | 13.022522 | JetBlue Airways |
| 8 | VX | 12.869421 | Virgin America |
| 9 | OO | 12.586207 | SkyWest Airlines Inc. |
| 10 | UA | 12.106073 | United Air Lines Inc. |
| 11 | MQ | 10.552041 | Envoy Air |
| 12 | DL | 9.264505 | Delta Air Lines Inc. |
| 13 | AA | 8.586016 | American Airlines Inc. |
| 14 | AS | 5.804775 | Alaska Airlines Inc. |
| 15 | HA | 4.900585 | Hawaiian Airlines Inc. |
| 16 | US | 3.782418 | US Airways Inc. |

Thank you

Question?